

# Bhuwaneshwar Sagar Chintala

Houston, TX | +1 713-815-1609 | bsagar0706@gmail.com | LinkedIn | GitHub

## SUMMARY

Data Engineer with 4+ years of experience building scalable cloud data platforms across healthcare and financial services. Specialized in cloud-native (AWS & GCP) architectures processing millions of daily events, dimensional modeling with dbt, and data quality frameworks ensuring SOX and HIPAA compliance for analytics and ML workloads.

## SKILLS

**Programming & Scripting:** Python, SQL (Advanced), T-SQL, PL/SQL, R, Shell Scripting, Linux/Unix, Java

**Cloud & Data Platforms:** AWS (S3, Glue, Redshift, Lambda, Step Functions, Kinesis, EMR), GCP (BigQuery, Dataflow, Pub/Sub, Dataproc, Cloud Functions), Snowflake, Databricks

**Data Engineering:** Apache Spark, Apache Kafka, Apache Flink, REST API Integration, dbt, Airflow, Dagster, ETL/ELT

**Databases:** PostgreSQL, MySQL, MongoDB, DynamoDB

**Data Modeling:** Dimensional Modeling, Star Schema, Snowflake Schema

**DevOps & IaC:** Terraform, Docker, Kubernetes, CI/CD, Git, GitHub Actions, GitLab CI

**Data Quality & Analytics:** Great Expectations, dbt Tests, Tableau, Power BI, CloudWatch

## WORK EXPERIENCE

### Ring Therapeutics — Data Engineer (Contract)

New Jersey, USA | Dec 2024 – Present

- Built a HIPAA-compliant AWS data pipeline using medallion architecture with S3, Glue, and Redshift to process de-identified healthcare records, ensuring robust RBAC and PHI encryption.
- Developed event-driven ingestion framework using Step Functions and Lambda to process 10K+ daily genomic records reducing processing costs from \$4.8K to \$1.7K via serverless orchestration.
- Engineered real-time Kafka + Flink pipelines with stateful windowing, achieving high reliability and meeting strict freshness SLAs monitored via CloudWatch.
- Collaborated with stakeholders to establish data quality and validation frameworks using Great Expectations integrated into CI/CD, implementing automated SNS alerting that identified 95% of data issues in pre-production.

### Ford Motor Company — Associate Data Engineer

Chennai, India | Jan 2022 – Dec 2023

- Built BigQuery warehouse with 5+ star schemas for 50+ analysts, accelerating query performance by 60% through optimized partitioning and clustering.
- Developed 30+ dbt models for connected vehicle analytics with incremental materialization and SCD Type-2 tracking, significantly improving data reliability through dbt tests and Dataplex profiling.
- Pioneered data migration from Teradata to BigQuery using Airflow and Dataproc, completing the project 2 weeks early through optimized broadcast joins and Spark caching.
- Provisioned GCP infrastructure via Terraform with Cloud KMS and IAM policies, streamlining deployment workflows while enforcing strict SOX and PCI-DSS compliance for sensitive credit data.
- Orchestrated 10+ Airflow DAGs managing end-to-end ETL workflows, data pulling from REST APIs with self-healing retry logic, custom dependencies and improved production pipeline uptime from 96.2% to 99.7%.

### Highradius Technologies — Machine Learning Engineer

Hyderabad, India | Jun 2020 – Jan 2022

- Built data pipelines processing 100K+ transactions from MySQL/PostgreSQL, utilizing automated validation and feature engineering to enhance model training accuracy.
- Developed NLP pipeline using spaCy/NLTK to process 50K+ documents, reducing manual review overhead and improving routing accuracy for dispute categorization.
- Deployed ensemble models (XGBoost/Random Forest) via Docker, enabling the recovery of \$2M+ in outstanding receivables and reducing DSO by 8 days.
- Created 5+ Tableau dashboards visualizing account prioritization, enabling leadership to accelerate decision-making through real-time collection performance metrics.

## EDUCATION

### University of Houston — Master of Engineering in Data Science (CGPA: 3.51 / 4.0)

May 2025

## CERTIFICATIONS

AWS Certified Data Engineer - Associate (Expires Aug 2028) | dbt Fundamentals - dbt Labs